

# Go-Shorty: A Comprehensive Technical Deep Dive

Sumanth

August 24, 2025

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Core Concepts &amp; Architecture</b>	<b>2</b>
2.1	Architectural Breakdown . . . . .	2
2.2	Key Data Flows . . . . .	3
2.2.1	Flow 1: Creating a Link (Synchronous) . . . . .	3
2.2.2	Flow 2: Redirecting a Link & Tracking a Click (Asynchronous) . . . . .	3
<b>3</b>	<b>Technology Deep Dive</b>	<b>4</b>
3.1	Go (golang) . . . . .	4
3.2	Chi (Router) . . . . .	4
3.3	PostgreSQL (Database) . . . . .	4
3.4	sqlc (Type-Safe SQL) . . . . .	4
3.5	Goose (Migrations) . . . . .	5
3.6	Redis (Cache & Queue) . . . . .	5
3.7	Docker & Docker Compose . . . . .	5
3.8	GitHub Actions (Continuous Integration) . . . . .	5
<b>4</b>	<b>Project Structure</b>	<b>5</b>

## 1 Introduction

Go-Shorty is a full-featured URL shortener engineered to be robust, scalable, and production-ready. This document provides a comprehensive, in-depth exploration of its architecture, the rationale behind the technology stack, and the specific design patterns implemented throughout its development.

The primary goal of this project is to serve as a real-world, canonical example of building a modern web application in Go. It moves beyond a simple "Hello, World" to tackle challenges inherent in production systems, including database management, asynchronous processing, stateful authentication, caching strategies, and containerized deployment. It solves the core problem of converting long, unwieldy URLs into manageable short links while providing essential features like custom aliases, secure user accounts, and detailed click analytics.

## 2 Core Concepts & Architecture

The application is built on a **decoupled, service-oriented architecture**. This is a foundational design choice that prioritizes performance and scalability. Instead of a single, monolithic application, the system is composed of two primary, independent applications—a **Web Server** and a **Click Worker**—that communicate asynchronously through a message queue.

This separation is the key to a responsive user experience. The user-facing **Web Server** is optimized for one thing: handling HTTP requests as quickly as possible. Resource-intensive, non-critical tasks, like writing analytics data to the database, are offloaded to the **Click Worker**. This ensures that spikes in traffic or slow database writes for analytics never impact the speed of user-facing operations like link redirection.

Let the set of system components be defined as  $S$ :


$$S = \{\text{Browser, Go Server, Go Worker, Redis, PostgreSQL}\}$$

### 2.1 Architectural Breakdown

Here is the final architecture, showing how the containerized services interact within the Docker network. This model represents a complete, self-contained environment that is both reproducible locally and deployable to the cloud.

**Flow 1: How a New Link is Created**

1. you submit the form in your browser. The browser sends a JSON API request (`POST /api/links`) to the Go Web Server.
2. The Go Web Server receives the request and saves the new link directly into the PostgreSQL Database.
3. The Go Web Server sends a confirmation back to your Browser, and the JavaScript on the page adds the new link to your dashboard.

**Flow 2: How a Click is Tracked**

1. A user clicks a short link (`/my-alias`). Their browser makes a request to the Go Web Server.
2. The Go Web Server first checks the Redis Cache. This is extremely fast. If the link isn't in the cache, it gets it from PostgreSQL and saves it to the cache for next time.
3. The server immediately puts a small "a click happened" message into the Redis Queue. This is a "fire-and-forget" action; the server doesn't wait for a response.
4. The server instantly sends a redirect back to the user's Browser, sending them to the final destination. This is the end of the user's interaction. It's very fast.
- Meanwhile, in the background:
5. The Go Click Worker, which is always listening, sees the new message in the Redis Queue and picks it up.
6. The Worker processes the message.
7. The Worker saves the detailed click information into the `clicks` table in the PostgreSQL Database.

## 2.2 Key Data Flows

### 2.2.1 Flow 1: Creating a Link (Synchronous)

This is a straightforward, synchronous, and atomic operation that prioritizes data consistency.

1. The user's browser, running our JavaScript front end, initiates a `POST` request with a JSON payload to the `/api/links` endpoint on the **Go Server**.
2. The server's handler receives the request, decodes the JSON into a Go struct, and passes it to the **Service Layer**.
3. The service layer performs business logic validation (e.g., checking for reserved aliases, validating the URL format).
4. It then calls the **Repository Layer**, which executes a SQL `INSERT` statement within a transaction to save the new link in the **PostgreSQL** database. The `UNIQUE` constraint on the `alias` column ensures data integrity at the database level.
5. Upon successful insertion, the database returns the newly created link record. The server sends a `201 Created` status code and the new link data back to the browser, which then updates the UI.

### 2.2.2 Flow 2: Redirecting a Link & Tracking a Click (Asynchronous)

This process is heavily optimized for speed, offloading all non-essential work to the background.

1. A user visits a short link (e.g., `/fJ7xP2`). Their browser sends a `GET` request to the **Go Server**.

2. The server's redirect handler first attempts to retrieve the destination URL from **Redis**. The cache lookup logic can be represented as:

$$\text{Destination} = \begin{cases} \text{Redis.GET}(\text{alias}) & \text{if Redis.EXISTS}(\text{alias}) \\ \text{PostgreSQL.SELECT}(\text{alias}) & \text{otherwise} \end{cases}$$

3. Concurrently, the server publishes a "click event" message to a Redis Stream. This is a "fire-and-forget" operation; the server adds the message to the stream and immediately moves on without waiting for any confirmation.
4. The server instantly sends an HTTP 302 Found redirect response to the browser, sending the user to their final destination. The user's interaction is now complete.
5. Meanwhile, the **Go Worker**, which has a blocking connection to the Redis Stream, receives the click event message.
6. The worker decodes the message and executes an INSERT statement, saving the detailed click information to the `clicks` table in **PostgreSQL**. This operation happens entirely independently of the user's redirect.

### 3 Technology Deep Dive

Each piece of technology was chosen for its specific strengths in building reliable, high-performance systems.

#### 3.1 Go (golang)

**Why?:** Go is the ideal language for this project. As a statically typed, compiled language, it produces small, high-performance binaries with no external runtime dependencies, which is perfect for containerization. Its built-in support for concurrency via **goroutines** and channels is fundamental to our architecture, allowing us to fire off the Redis publish event in a separate goroutine without blocking the main HTTP response thread.

#### 3.2 Chi (Router)

**Why?:** While Go's standard `http.ServeMux` is capable, **chi** provides a more expressive and powerful API for building RESTful services. It is lightweight and offers critical features like a powerful middleware system, URL parameter extraction, and route grouping.

#### 3.3 PostgreSQL (Database)

**Why?:** For an application requiring strong data integrity, PostgreSQL is the superior choice. Its **ACID compliance** guarantees that transactions are processed reliably. Relational features like UNIQUE constraints and foreign key relationships with `ON DELETE CASCADE` are essential for maintaining a clean and consistent state.

#### 3.4 sqlc (Type-Safe SQL)

**Why?:** `sqlc` is a transformative tool that generates fully type-safe Go code directly from your raw SQL queries. This eliminates an entire class of bugs by catching SQL syntax errors and data-mapping mismatches at compile time, not at runtime.

### 3.5 Goose (Migrations)

**Why?:** Managing database schema changes over time is critical. **Goose** provides a programmatic and version-controlled way to manage these changes, ensuring that every environment is always in a consistent, known state.

### 3.6 Redis (Cache & Queue)

**Why?:** Redis is the industry standard for high-speed, in-memory data operations. We use it for two distinct roles:

- **Caching:** Simple **GET/SET** commands with an expiration time (TTL) to cache link destinations.
- **Queuing:** We use **Redis Streams**, a persistent message queue. The server adds a click event with **XADD**, and the worker consumes it with a blocking **XREADGROUP** command.

### 3.7 Docker & Docker Compose

**Why?:** Docker packages our applications into isolated, reproducible images. Docker Compose orchestrates the entire multi-container environment (server, worker, database, cache) with a single command.

### 3.8 GitHub Actions (Continuous Integration)

**Why?:** CI is a fundamental practice of modern software development. By automatically building our code on every push, we ensure that no breaking changes are integrated into the main branch.

## 4 Project Structure

The project follows the standard Go project layout to ensure a clean separation of concerns.

- **cmd/:** Main application entry points (**server/main.go**, **worker/main.go**).
- **internal/:** All the core application code (config, handlers, middleware, repositories, services).
- **migrations/:** SQL migration files managed by **Goose**.
- **pkg/:** Reusable utility packages.
- **static/:** All the front-end files (HTML, CSS, JavaScript).
- **Dockerfile:** Recipe for building the application's Docker image.
- **docker-compose.yaml:** Defines how to run the entire application stack.
- **go.mod / go.sum:** Go's dependency management files.